

A classification of well-behaved graph clustering schemes

Vilhelm Agdur

28 September 2023

1 Introduction

The problem of inferring community structure from a graph is one that appears in many different guises in different research areas. It is useful for everything from assigning research papers to subfields of physics [CR10], assessing the development of nano-medicine by studying patent cocitation networks [BAB12], studying brain function via the community structure of the connectome [Bet23], understanding how different bird species spread various plant seeds [Cos+16], to studying hate speech on Twitter [Evk+21].

Despite, or perhaps because of, the widespread nature of this problem, there is no universally agreed on definition of what is actually meant by community structure. This leads to varying methods being used to solve the problem, that do not necessarily give the same results. One of the most popular definitions in practice is the modularity of a partition, introduced by Newman and Girvan [NG04] – it is the method used in all the examples given in the previous paragraph. There are however many other methods, such as the flow-based infomap method motivated by information theory [RAB09], novel graph neural network methods [Koj+23], and many more [Jav+18].

One thing nearly all of these methods have in common is that they are in a sense black boxes – they do not give you any way of answering the question of *why* two vertices were put in the same part. They are also generally hard to study in a rigorous mathematical way, resulting in few theoretical guarantees for their behaviour. While modularity has proven itself very useful in practice, it is in fact guaranteed to behave in undesirable ways in certain settings.

In particular, modularity suffers from a “resolution limit”, wherein it cannot “see” things happening in subgraphs that are small compared to the total graph, resulting in very unexpected behaviour [FB07]. In Figure 1 we see one illustration of this phenomenon – one would hope that what happens in one connected component cannot affect the other, but the resolution limit results in this undesirable behaviour.

Another issue with studying modularity optimization from a theoretical perspective is the fact that computing it is likely NP-hard.¹ This means that when using modularity maximization in practice, one always needs a heuristic algorithm for finding good partitions, such as the Louvain [Blo+08] or Leiden [TWV19] algorithms, which themselves do not have many known theoretical guarantees connecting them to modularity. So one is left having to either study modularity itself, and hoping that the results say something about the results

¹The paper originally proving this, [Bra+08], has an error in a lemma that, at the time of writing, has not yet been fixed. It does however seem likely that the result holds, and it is certainly still true that we know of no tractable exact algorithm even if one could exist.



(a) A graph G , with vertices coloured according to the modularity-maximizing partition (b) G with a single edge added, highlighted in red, and its new optimal partition

Figure 1: An illustration of the resolution limit issue with modularity.

of these algorithms, or studying the algorithms directly, and not directly saying anything about modularity.

These issues with the varying optimization based approaches to clustering lead us to investigating a more “axiomatic” approach to the problem – that is, we want to specify a few properties that the clustering method has to have, and then seeing what methods there are that satisfy these properties. Other than of course giving us guarantees on the behaviour of our methods, this approach also gives us much more straightforward algorithms which run in polynomial time, since we are no longer doing optimization, but in some sense just applying a decision rule for when vertices should be in the same part or not.

Our particular approach to this is motivated by the results of Carlsson and Mémoli [CM13] for the analogous problem of clustering finite metric spaces. Like in their setting, there are three properties that are of interest:

1. *Functoriality* is essentially a form of monotonicity under addition of edges or vertices to the graph. Adding a new edge or new vertex should only ever make communities more tightly connected, it should never split them.
2. *Excisiveness* is a very strict version of requiring that there be no resolution limit – instead we require that the method be idempotent on the parts it has found. That is, if we take some graph G , cluster it, and then look at the subgraph induced by any part, that graph has to be clustered into just a single large component, so we never discover new features at smaller scales.
3. *Representability* is a type of explainability of the clustering method. A representable clustering method F_Ω is determined by a set Ω of simple graphs, and the rule that two vertices of a graph G will be in the same part if their neighbourhood looks like something in Ω . This also gives us a simple algorithm to actually compute the partitioning of a graph, that runs in polynomial time for each fixed Ω .

Our main result is that for functorial clustering schemes, the properties of excisiveness and representability are equivalent. We also give a structural result for when two representable clustering schemes are equal, and use this to prove that there exist representable clustering

schemes that are not finitely representable. Finally, we prove that on any class with bounded expansion, we can, for any finite representing set Ω , in fact compute the part to which a vertex belongs in time $O(n^2) + O(w)$, where w is the number of subgraphs isomorphic to something in the representing set.

We also choose to extend our definition of a representable clustering scheme into the setting of hierarchical clustering schemes. Here, instead of just assigning a single clustering to each graph, we give one clustering for each “scale”, giving a spectrum of clusterings of the graph from coarsest to finest. In this setting, we find that the connection to excisiveness no longer works out.

2 Definitions of our terms

Definition 1. The category of simple graphs $\mathcal{G}$ has as its objects simple graphs, which we consider to be a finite set V equipped with a reflexive and symmetric binary relation E – that is, our simple graphs are finite, have no double edges, and have loops at every vertex². A morphism from $H = (V, E)$ to $G = (V', E')$ is an injective set function $f : V \rightarrow V'$ such that whenever uEv , we also have $f(u)E'f(v)$. That is, a morphism from H to G is a way of realizing H as a subgraph of G .

Definition 2. The category of partitioned sets $\mathcal{P}$ has as its objects sets V with an equivalence relation $\sim$. A morphism in $\mathcal{P}$ from $(V, \sim)$ to $(W, \approx)$ is a set function $f : V \rightarrow W$ such that whenever $v \sim w$, $f(v) \approx f(w)$. So in the case where $V = W$, $f = \text{Id}$ is a morphism if $\approx$ is a coarsening of $\sim$.

Definition 3. A clustering scheme is a function³ $\mathfrak{C}$ that sends simple graphs to partitioned sets, with the same underlying set. We say that it is *functorial* if it is a functor from $\mathcal{G}$ to $\mathcal{P}$.

Definition 4. A clustering scheme $\mathfrak{C}$ is *excisive* if, for all graphs $G = (V, E)$, it holds for every equivalence class A of $\mathfrak{C}(G)$ that $\mathfrak{C}(G|_A) = (A, A^2)$. That is, if we take a graph, cluster it, pick out one of the parts, and cluster that part alone, that part will be clustered into a single part.

3 Representability

Definition 5. Given a not necessarily finite set Ω of simple graphs, we define the endofunctor F_Ω on $\mathcal{G}$ as follows: For any graph $G = (V, E)$, $F_\Omega(G)$ is a graph on the same vertex set, and there is an edge between u and v in $F_\Omega(G)$ whenever there exists some subgraph of G containing both u and v which is isomorphic to some $\omega \in \Omega$.

Equivalently, there is an edge uv in $F_\Omega(G)$ whenever there exists an $\omega \in \Omega$ and a morphism $f : \omega \rightarrow G$ such that $f(\omega) \ni u, v$.

On morphisms, F_Ω is always just the identity.

Lemma 6. *For any representing set Ω of simple graphs, F_Ω is indeed an endofunctor on $\mathcal{G}$.*

²The final property here may seem a bit unusual, but it simplifies our proofs in some places, without really affecting much.

³Technically, of course, neither $\mathcal{G}$ nor $\mathcal{P}$ are small categories, so the word “function” should be read as “class function” here, not in the usual sense of the word. We sweep this issue under the rug here, because it will never affect anything.

Proof. That it always gives a simple graph and that it respects identity morphisms and composition of morphisms is easy to see. The only thing we actually need to check is that if $f : H \rightarrow G$ is a morphism, then $F_\Omega(f) = f : F_\Omega(H) \rightarrow F_\Omega(G)$ is also a morphism.

Unwrapping this statement, what we need is that, for any morphism $f : H \rightarrow G$, whenever uv is an edge in $F_\Omega(H)$, $f(u)f(v)$ is an edge in $F_\Omega(G)$. Now, that uv is an edge in $F_\Omega(H)$ means that there is an $\omega \in \Omega$ and $f_\omega : \omega \rightarrow H$ such that $f(\omega) \ni u, v$.

But consider what happens if we just compose f_ω and f – this will be a morphism from ω into G , and trivially we must have that $f(f_\omega(\omega)) \ni f(u), f(v)$, and so there must be an edge $f(u)f(v)$ in $F_\Omega(G)$, by definition of F_Ω . $\square$

What is going on in the proof of Lemma 6 is much clearer if we just think about it in terms of things being subgraphs of each other, forgetting the data of *how* exactly they are subgraphs. Then the statement that F_Ω is a functor boils down to the statement that if ω is a subgraph of H and H is a subgraph of G , then certainly ω is a subgraph of G , and so $F_\Omega(G)$ must have all the edges $F_\Omega(H)$ has, since edges in $F_\Omega(G)$ represent the presence of a certain subgraph in G .

Definition 7. The functor $\Pi : \mathcal{G} \rightarrow \mathcal{P}$ which sends a graph to its partition into connected components is called the connected components functor. Equivalently, we can think of it as taking the equivalence relation hull of the binary relation E defining the edges.

Definition 8. A clustering scheme $\mathfrak{C}$ is *representable* if we can write it as $F_\Omega \circ \Pi$ for some representable endofunctor F_Ω . Notice that this means that all representable clustering schemes are functorial, being the composition of two functors.

One somewhat surprising fact is that there are actually representable clustering schemes that are not finitely representable.

Theorem 9. *There exists a representable endofunctor F_Ω which is not finitely representable, that is, it is not equal to F_Γ for any finite set of simple graphs Γ .*

In order to give a proof of this, we first need a structural result about how these endofunctors behave.

Remark 10. For any $\omega \in \Omega$, it is trivial to see that $F_\Omega(\omega)$ must be a clique, since we can just take the identity morphism on ω to show that all edges must exist.

In fact, which graphs are sent to cliques by F_Ω essentially determines what it does as a functor. We will see this even more precisely in the next section, but for now this lemma will suffice for our needs:

Lemma 11. *For any set Ω of simple graphs and any simple graph $G = (V, E)$, it holds that $F_{\Omega \cup \{G\}} = F_\Omega$ if and only if $F_\Omega(G) = (V, V^2)$, that is, if G is sent to a clique on its vertices by F_Ω .*

Proof. Suppose $F_{\Omega \cup \{G\}} = F_\Omega$. That $F_{\Omega \cup \{G\}}(G) = (V, V^2)$ is immediate by Remark 10, and so since $F_{\Omega \cup \{G\}} = F_\Omega$, we also have $F_\Omega(G) = (V, V^2)$, as desired.

Now suppose $F_\Omega(G) = (V, V^2)$, and let H be some arbitrary simple graph. We wish to show that $F_{\Omega \cup \{G\}}(H) = F_\Omega(H)$.

That edges of $F_\Omega(H)$ must also be edges of $F_{\Omega \cup \{G\}}(H)$ is obvious, so all we need to show is that edges of $F_{\Omega \cup \{G\}}(H)$ are also edges of $F_\Omega(H)$. So suppose uv is an edge of $F_{\Omega \cup \{G\}}(H)$ – that this is so must be because there is some $\omega \in \Omega \cup \{G\}$ and an $f : \omega \rightarrow H$ such that $f(\omega) \ni u, v$.




Figure 2: The family of graphs $\{\omega_i\}_{i=0}^3$.

If this ω is not G , then we are done. If it is G , we recall that since $F_\Omega(G)$ is a clique, there must be an edge $f^{-1}(u)f^{-1}(v)$ in it. Therefore there must exist some $\omega \in \Omega$ and some $g : \omega \rightarrow G$ such that $f(\omega) \ni f^{-1}(u), f^{-1}(v)$. But then the composition $g \circ f : \omega \rightarrow H$ will be a morphism from an $\omega \in \Omega$ into H whose image contains u and v , proving that uv is an edge of $F_\Omega(H)$. $\square$

With this result in hand, we can now give a proof of the existence of a representable but not finitely representable endofunctor.

Proof of Theorem 9. Let, for each i , ω_i consist of a triangle with a tail of i vertices, as shown in Figure ??, and let $\Omega = \{\omega_i \mid i \in \mathbb{N}\}$. We claim that F_Ω is not equal to F_Γ for any finite set of simple graphs Γ .

To prove this, let us first define

$$\bar{\Omega} = \{G \in \mathcal{G} \mid F_\Omega(G) = (V_G, V_G^2)\}$$

as the set of all things which F_Ω sends to cliques. By Lemma 11, $F_\Omega = F_{\bar{\Omega}}$, and $\bar{\Omega}$ is the greatest set with this property.

A little bit of reflection will make it clear that $\bar{\Omega}$ is precisely the set of connected graphs which contain a triangle. Particularly, for any edge in such a graph, we can map the tail end of some ω_i onto that edge, the rest of the tail onto a path to the triangle in the graph, and the triangle of the ω_i onto the triangle of the graph. That F_Ω will send any triangle-free graph to a graph with no edges is clear, since there is no morphism from any triangle-containing graph such as ω_i into a triangle-free graph. Likewise, that it cannot send a disconnected graph to a clique is obvious.

Now suppose for contradiction that Γ is a finite set of graphs such that $F_\Gamma = F_{\bar{\Omega}}$. Lemma 11 implies that we must have $\Gamma \subset \bar{\Omega}$, since if there were a $\gamma \in \Gamma \setminus \bar{\Omega}$ it'd be sent to a clique by F_Γ , and thus also by $F_{\bar{\Omega}}$, since we've assumed they are equal – but we chose $\bar{\Omega}$ so that it already contains everything sent to a clique by $F_{\bar{\Omega}}$.

So, let

$$\delta = \max_{\gamma \in \Gamma} \max_{u \in V_\gamma} \min_{\substack{v \in V_\gamma \\ v \text{ is in a triangle}}} d(u, v).$$

That is, δ is the furthest any vertex in any graph in Γ is from a triangle. We claim that F_Γ will not send $\omega_{\delta+1}$ to a clique, proving it is not equal to $F_{\bar{\Omega}}$, giving our contradiction.

That this is so is actually easy to see – if $\omega_{\delta+1}$ were sent to a clique, that must in particular mean there is an edge in $F_\Gamma(\omega_{\delta+1})$ between the tail vertex and its neighbour. So there is some morphism from some $\gamma \in \Gamma$ that hits both vertices. However, such a morphism must

also map the triangle in γ onto the triangle in $\omega_{\delta+1}$, since that is the only place the triangle can go. Then, however, it cannot also reach to the tail – the tail is further away from the triangle than any vertex in any γ is from the triangle. So no such morphism can exist, and we are done. $\square$

Remark 12. The construction in our proof of course works equally as well replacing the triangle by any connected graph that is not a line. So we have found a whole lot of examples of classes of graphs that give representable but not finitely representable endofunctors.

The question of which natural classes of graphs are representable or finitely representable may be interesting. The connected graphs, for example, are of course represented by just a K_2 . The example we gave in the proof easily generalizes to showing that the graphs of girth at most k are representable but not finitely representable, by extending from triangles with tails to all cycles of length at most k with tails.

The class of non-planar graphs is also representable, with itself as a representation, since there can of course be no morphism from a non-planar graph into a planar graph. Whether there is a smaller representing set, or even a finite one, is a mildly interesting question we leave unanswered.

4 Proof of equivalence of representability and excisiveness

We are able to exactly characterize the relationship between the three concepts of functoriality, representability, and excisiveness for clustering schemes.

Theorem 13. *A representable clustering scheme is excisive. An excisive and functorial clustering scheme is representable. No other implications hold between these concepts.*

We divide the proof of Theorem 13 into two lemmas and a collection of examples.

Lemma 14. *A representable clustering scheme is excisive.*

Proof. Let Ω be some set of simple graphs, and G be some simple graph. Assume the equivalence classes of $(F_\Omega \circ \Pi)(G)$ are $\{A_1, A_2, \dots, A_k\}$, and let $H = G|_{A_1}$ be the induced subgraph of G on A_1 . What we need to show is that $(F_\Omega \circ \Pi)(H) = (A_1, A_1^2)$, that is, that every vertex is equivalent in the clustering of H .

It is obvious that this is equivalent to showing that $F_\Omega(H)$ is a connected graph. We will in fact show something stronger – that $F_\Omega(H) = F_\Omega(G_{A_1}) = (F_\Omega(G))|_{A_1}$. Since A_1 is by assumption a connected component of $F_\Omega(G)$, this will show what we want.

First, let $f : H \hookrightarrow G$ be the subgraph inclusion morphism of H into G . By functoriality of F_Ω , f is also a morphism from $F_\Omega(H)$ into $F_\Omega(G)$, showing that $F_\Omega(H)$ is a subgraph of $F_\Omega(G)$, so every edge of $F_\Omega(H)$ is an edge of $F_\Omega(G)$.

Now, let uv be an arbitrary edge of $F_\Omega(G)|_{A_1}$. We wish to show this is already an edge of $F_\Omega(H)$. That it is an edge of $F_\Omega(G)$ means there exists some $\omega \in \Omega$ and a morphism $f_\omega : \omega \rightarrow G$ such that $\text{im}(f_\omega) \ni u, v$.

We claim that in fact $\text{im}(f_\omega) \subset A_1$, so that f_ω restricts to a morphism into H . If we can show this, this will show that there is an edge uv also in $F_\Omega(H)$, and we will be done.

So let w be an arbitrary vertex in $\text{im}(f_\omega)$. Now f_ω is also a morphism from something in Ω that hits both w and v , so there is an edge vw in $F_\Omega(G)$. But this of course means that

w and v are in the same connected component – and this connected component is precisely A_1 . So $w \in A_1$, and we are done. $\square$

Lemma 15. *An excisive and functorial clustering scheme is representable.*

Proof. Let $\mathfrak{C}$ be some excisive and functorial clustering scheme, and let

$$\Omega = \{G \in \mathcal{G} \mid \mathfrak{C}(G) = (V_G, V_G^2)\},$$

that is, Ω is the set of all graphs clustered by $\mathfrak{C}$ into a single part. We will show that in fact $\mathfrak{C} = F_\Omega$.

So, let G be some graph. We need to show two things:

1. Whenever u and v are clustered in the same part by $\mathfrak{C}$, they are in the same connected component of $F_\Omega(G)$.
2. Whenever uv is an edge of $F_\Omega(G)$, u and v are clustered in the same part by $\mathfrak{C}$. This of course implies that whenever u and v are in the same connected component, they are also clustered in the same part by $\mathfrak{C}$, which is what we really need.

We start with the first direction, letting u and v be two vertices that are sent to the same part by $\mathfrak{C}$ – call this part A , and let $H = G|_A$. By excisiveness of $\mathfrak{C}$, we have that $\mathfrak{C}$ sends H to a single part, so $H \in \Omega$. Therefore, the inclusion morphism $f : H \hookrightarrow G$ is a morphism from something in Ω that hits both u and v , showing there is an edge uv in $F_\Omega(G)$, as desired.

In the second direction, assume that uv is an edge of $F_\Omega(G)$. Thus, there is some $\omega \in \Omega$ and $f_\omega : \omega \rightarrow G$ such that $\text{im}(f_\omega) \ni u, v$. By functoriality of $\mathfrak{C}$, f will also be a morphism from $\mathfrak{C}(\omega)$ into $\mathfrak{C}(G)$ – which by how we defined morphisms in the category of partitioned sets means that any two vertices equivalent in $\mathfrak{C}(\omega)$ must also be equivalent in $\mathfrak{C}(G)$. However, we defined Ω as precisely the set of things such that $\mathfrak{C}(\omega)$ has only a single part – so everything is equivalent in $\mathfrak{C}(\omega)$ and in particular $u \sim_{\mathfrak{C}(\omega)} v$, and so they are also equivalent in $\mathfrak{C}(G)$, as desired. $\square$

Finally, let us think about what we need to show to demonstrate that there are no other implications between these concepts. An implication between these concepts is logically the same thing as a proof that some combination of them is impossible – so what we will do is just to draw a two-by-two-by-two table of all combinations, and fill in which cells our theorems already rule out, and give examples of clustering schemes for every other cell.

Functorial		
Not functorial	Excisive	Not Excisive
Representable	Yes, Ex. 16	No, by Lm. 14
	No, by Def. 8	No, by Def. 8 (or by Lm. 14)
Not Representable	No, by Lm. 15	Yes, Ex. 18
	Yes, Ex. 17	Yes, Ex. 19

Table 1: Table of all the possible combinations of properties, with examples or proofs of impossibility.

Example 16. The connected components functor Π is representable and excisive – we can take its representation to be $\Omega = \{K_2\}$.

Example 17. For a scheme that is excisive but not functorial (and thus not representable), consider the scheme that partitions all graphs into a single part except K_2 , which it partitions into two parts.

Example 18. For a scheme which is functorial but neither excisive nor representable, cluster all connected components on more than two vertices as well as all isolated vertices as their own parts.

If there is exactly one component on two vertices, i.e. an isolated edge, cluster it as two parts, otherwise cluster all isolated edges as single parts.

Example 19. Finally, for a scheme which has none of the properties, cluster all graphs as a single part except for K_2 , which we cluster as two parts, and the disjoint union of two copies of K_2 , which we cluster as two parts – that is, each of the copies is its own part.

Remark 20. It is possible to carry out this entire argument, showing equivalence between representability and excisiveness, also in the setting where we do not require morphisms in the category of graphs to be injective. The arguments only need minor modifications to deal with this.

However, it turns out that this setting is much less interesting: There are in fact only three different representable clustering schemes in this setting. They are, respectively,

1. the scheme that always sends every vertex to its own part, represented by the empty set or by K_1 ,
2. the connected components functor, represented by any set of connected graphs containing at least one graph on at least two vertices,
3. the scheme that always sends all vertices to the same part, represented by any set of graphs containing a disconnected graph.

5 Computational complexity of representable clustering schemes

As stated in the introduction, one benefit of this approach to clustering is that it yields tractable algorithms for exactly computing the partitioning. In this section, we give a proof of this.

The most naïve possible algorithm given a fixed set Ω of representing graphs and an n -vertex graph G is to, for each $\omega \in \Omega$ and each $H \in \binom{V(G)}{|\omega|}$, that is, each $|\omega|$ -sized subset of the vertices of G , check whether $G|_H$ contains a subgraph isomorphic to ω . It is easy to see that this will give an algorithm that is polynomial in n , but with an exponent and constant depending on Ω .

What is more interesting is that we can get a far more efficient algorithm on a large class of graphs, using the following lemma from [ND12, p. 407]:

Lemma 21. *Let $\mathcal{C}$ be a class with bounded expansion and let H be a fixed graph. Then there exists a linear time algorithm which computes, from a pair (G, S) formed by a graph $G \in \mathcal{C}$ and a subset S of vertices of G , the number of isomorphs of H in G that include some vertex*

in S . There also exists an algorithm running in time $O(n) + O(k)$ listing all such isomorphs, where k denotes the number of isomorphs.

Theorem 22. Let Ω be any finite set of simple graphs, and F_Ω the representable clustering scheme represented by it. There exists an algorithm that given any graph G and vertex v of G computes the part containing v of $F_\Omega(G)$ running in time $O(n^2) + O(w)$, where w is the number of subgraphs of G isomorphic to a graph in Ω .

Proof. It is clear that if we can compute the part containing v in $F_{\{\omega\}}(G)$ for each $\omega \in \Omega$, then we can just take the union of these parts to get v 's part in $F_\Omega(G)$ – and if we can do this for each ω in time $O(n^2) + O(w_\omega)$, with w_ω the number of copies of ω in G^4 , it will sum up to the running time we give in the statement.

So, let us restrict to the case of $\Omega = \{\omega\}$. We will of course use the algorithm from Lemma 21, and we do so in the most straightforward way.

Algorithm 1 Computation of part of v in G

```

 $P_0 \leftarrow \{v\}, P_{-1} \leftarrow \emptyset$ 
for  $i = 1, 2, \dots, n$  do
  if  $P_{i-1} \setminus P_{i-2} \neq \emptyset$  then
    Run the algorithm of Lemma 21 with input  $(G, P_{i-1} \setminus P_{i-2})$ , getting a list of
    isomorphs  $\gamma_1, \dots, \gamma_k$ 
     $P_i \leftarrow P_{i-1} \cup \bigcup_{j=1}^k V(\gamma_j)$ 
  else
    return  $P_{i-1}$ 
  end if
end for
return  $P_n$ 

```

Each step in the loop takes time $O(n)$ plus the number of isomorphs found in the step, so since the loop can take at most n steps, the total runtime must be at most $O(n^2)$ plus the total number of isomorphs in the graph, as desired. $\square$

6 Representability for hierarchical clustering

Definition 23. A functorial hierarchical clustering scheme is a functor from the product category $\mathcal{G} \times \mathbb{R}$ into $\mathcal{P}$, where we make $\mathbb{R}^+ = [0, \infty)$ into a category by saying there is a morphism from a to b whenever $a \geq b$.

There is actually a natural extension of our definition of a representable clustering scheme also to this setting, but it requires that we set up some more machinery.

Definition 24. The category $\mathcal{G}^w$ of weighted simple graphs has as its objects simple graphs together with a function assigning a real-valued positive weight to each edge of the graph. A morphism from (V, E, w) to (V', E', w') is an injective set function $f : V \rightarrow V'$ such that whenever uEv , $f(u)E'f(v)$ and $w(uv) \leq w'(f(u)f(v))$.

⁴There is a standard notation for this, right, from extremal graph theory? I just don't remember it right now.

Definition 25. The scissors functor Σ from $\mathcal{G}^w \times \mathbb{R}^+$ into $\mathcal{G}$ takes in a weighted graph $G = (V, E, w)$ and a non-negative real number τ , removes all edges whose weight is below τ , and then forgets all the remaining weights, getting just a simple graph.

Lemma 26. *The scissors functor is actually a functor.*

Proof. Similarly to how we proved that the representable endofunctors are in fact functors, most of the required properties are obvious. There are only two things we need to check:

1. If f is a morphism of weighted graphs sending $G = (V, E, w)$ to $G' = (V', E', w')$, then for every $\tau \in \mathbb{R}^+$, f is also a morphism in $\mathcal{G}$ from $\Sigma(G, \tau)$ to $\Sigma(G', \tau)$,
2. and if $\alpha \geq \beta \geq 0$, then for any weighted graph $G = (V, E, w)$, the identity function $\text{Id} : V \rightarrow V$ is a morphism from $\Sigma(G, \alpha)$ to $\Sigma(G, \beta)$.

We start with the first item, and suppose that $f : G \rightarrow G'$ is a morphism of weighted graphs, and $\tau \geq 0$ some threshold. We wish to show that whenever uv is an edge of $\Sigma(G, \tau)$, $f(u)f(v)$ is an edge of $\Sigma(G', \tau)$. That uv is an edge of $\Sigma(G, \tau)$ means that uv is an edge of G with weight at least τ , and since f is a morphism of weighted graphs, $f(u)f(v)$ must be an edge of G' , and since such morphisms are weight non-decreasing, its weight must still be at least τ . So it is an edge of $\Sigma(G', \tau)$ as well.

The second item is even more trivial than the first – if we unwrap the definitions, all it says is that if we lower the threshold for when edges are removed, then we will not lose any edges. $\square$

Definition 27. A representable functor F_Ω from $\mathcal{G}$ to $\mathcal{G}^w$ is determined by a set Ω of pairs (ω, w) of simple graphs ω and positive weights w . For any graph G , there is an edge uv of $F_\Omega(G)$ if there is some subgraph of G containing u and v which is isomorphic to some graph in Ω . The weight of this edge is then given by

$$w(uv) = \sum_{(\omega, w) \in \Omega} \sum_{f \in \text{Hom}(\omega, G)} \frac{w}{|\text{Aut}(\omega)|} \mathbb{1}_{u, v \in \text{im}(f)}.$$

Lemma 28. *The representable functors from $\mathcal{G}$ to $\mathcal{G}^w$ are actually functors.*

Proof. As before, the only thing we need to explicitly check is that whenever $f : G \rightarrow G'$ is a morphism in $\mathcal{G}$, it is also a morphism from $F_\Omega(G)$ to $F_\Omega(G')$. So, in particular, assume that uv is an edge of $F_\Omega(G)$ with weight α – what we need to show is that $f(u)f(v)$ is an edge of $F_\Omega(G')$ with weight at least α .

That it is an edge at all is easy to see – that it is an edge in $F_\Omega(G)$ means there is some $\omega \in \Omega$ and a morphism $f_\omega : \omega \rightarrow G$ whose image contains both u and v . If we just compose this f_ω with f , we get a morphism from ω into G' whose image of course contains $f(u)$ and $f(v)$.

To show that its weight is at least α , we compute

$$\begin{aligned}
w_{F_\Omega(G')}(uv) &= \sum_{(\omega,w) \in \Omega} \sum_{f'_\omega \in \text{Hom}(\omega, G')} \frac{w}{|\text{Aut}(\omega)|} \mathbb{1}_{f(u), f(v) \in \text{im}(f'_\omega)} \\
&\geq \sum_{(\omega,w) \in \Omega} \sum_{\substack{f'_\omega \in \text{Hom}(\omega, G') \\ \exists f_\omega \in \text{Hom}(\omega, G): f'_\omega = f_\omega \circ f}} \frac{w}{|\text{Aut}(\omega)|} \mathbb{1}_{f(u), f(v) \in \text{im}(f'_\omega)} \\
&= \sum_{(\omega,w) \in \Omega} \sum_{f_\omega \in \text{Hom}(\omega, G)} \frac{w}{|\text{Aut}(\omega)|} \mathbb{1}_{f(u), f(v) \in \text{im}(f_\omega \circ f)} \\
&= \sum_{(\omega,w) \in \Omega} \sum_{f_\omega \in \text{Hom}(\omega, G)} \frac{w}{|\text{Aut}(\omega)|} \mathbb{1}_{u, v \in \text{im}(f_\omega)} = w_{F_\Omega(G)}(uv) = \alpha,
\end{aligned}$$

where the inequality is just that we sum over a subset of the summands, and the second equality follows from that our morphisms are injective. $\square$

Definition 29. A representable hierarchical clustering scheme is a functor from $\mathcal{G} \times \mathbb{R}^+$ into $\mathcal{P}$ that can be written as $(F_\Omega \times \text{Id}) \circ \Sigma \circ \Pi$ for some representable functor F_Ω from $\mathcal{G}$ to $\mathcal{G}^w$.

Remark 30. We can of course turn such a representable hierarchical clustering scheme into a non-hierarchical clustering functor by just fixing a threshold. However, unfortunately, it turns out that this will *not* in general result in an excisive, and thus also not a representable, functor!




Figure 3: An example of a graph showing that $\Omega = \{(K_3, 1), (C_4, 1)\}$ does not give an excisive clustering functor at threshold 1.5.

To see this, take as your representing set the triangle and the four-cycle, both with weight one, and consider what is happening in Figure 3. If we apply the scissors functor at threshold 1.5, we will keep only the edge between u and v , so they get clustered as one part. However, if we zoom in on just this part, and try to cluster this K_2 , it will get clustered as two parts. So we have seen that at this threshold, it is in fact not excisive.

References

- [BAB12] Ahmad Barirani, Bruno Agard, and Catherine Beaudry. “Discovering and assessing fields of expertise in nanomedicine: a patent co-citation network perspective”. In: *Scientometrics* 94.3 (Nov. 2012), pp. 1111–1136. DOI: 10.1007/s11192-012-0891-6. URL: <https://doi.org/10.1007/s11192-012-0891-6> (cit. on p. 1).

- [Bet23] Richard F. Betzel. “Chapter 7 - Community detection in network neuroscience”. In: *Connectome Analysis*. Ed. by Markus D Schirmer, Tomoki Arichi, and Ai Wern Chung. Academic Press, 2023, pp. 149–171. ISBN: 978-0-323-85280-7. DOI: <https://doi.org/10.1016/B978-0-323-85280-7.00016-6>. URL: <https://www.sciencedirect.com/science/article/pii/B9780323852807000166> (cit. on p. 1).
- [Blo+08] V. D. Blondel et al. “Fast unfolding of communities in large networks”. In: *Journal of Statistical Mechanics: Theory and Experiment* 2008.10 (2008), P10008 (cit. on p. 1).
- [Bra+08] U. Brandes et al. “On modularity clustering”. In: *Knowledge and Data Engineering, IEEE Transactions on* 20.2 (2008), pp. 172–188. DOI: 10.1109/TKDE.2007.190689 (cit. on p. 1).
- [CM13] Gunnar Carlsson and Facundo Mémoli. “Classifying clustering schemes”. In: *Foundations of Computational Mathematics* 13.2 (Jan. 2013), pp. 221–252. DOI: 10.1007/s10208-012-9141-9. URL: <https://doi.org/10.1007/s10208-012-9141-9> (cit. on p. 2).
- [Cos+16] José M. Costa et al. “Sampling completeness in seed dispersal networks: When enough is enough”. In: *Basic and Applied Ecology* 17.2 (2016), pp. 155–164. ISSN: 1439-1791. DOI: <https://doi.org/10.1016/j.baae.2015.09.008>. URL: <https://www.sciencedirect.com/science/article/pii/S1439179115001255> (cit. on p. 1).
- [CR10] P. Chen and S. Redner. “Community structure of the physical review citation network”. In: *Journal of Informetrics* 4.3 (2010), pp. 278–290. ISSN: 1751-1577. DOI: <https://doi.org/10.1016/j.joi.2010.01.001>. URL: <https://www.sciencedirect.com/science/article/pii/S1751157710000027> (cit. on p. 1).
- [Evk+21] Bojan Evkoski et al. “Evolution of topics and hate speech in retweet network communities”. In: *Applied Network Science* 6.1 (Dec. 2021). DOI: 10.1007/s41109-021-00439-7. URL: <https://doi.org/10.1007/s41109-021-00439-7> (cit. on p. 1).
- [FB07] Santo Fortunato and Marc Barthélemy. “Resolution limit in community detection”. In: *Proceedings of the National Academy of Sciences of the United States of America* 104.1 (Jan. 2007), pp. 36–41. DOI: 10.1073/pnas.0605965104. URL: <https://doi.org/10.1073/pnas.0605965104> (cit. on p. 1).
- [Jav+18] Muhammad Aqib Javed et al. “Community detection in networks: A multidisciplinary review”. In: *Journal of Network and Computer Applications* 108 (2018), pp. 87–111. ISSN: 1084-8045. DOI: <https://doi.org/10.1016/j.jnca.2018.02.011>. URL: <https://www.sciencedirect.com/science/article/pii/S1084804518300560> (cit. on p. 1).
- [Koj+23] Sadamori Kojaku et al. *Network community detection via neural embeddings*. 2023. arXiv: 2306.13400 [physics.soc-ph] (cit. on p. 1).
- [ND12] Jaroslav Nešetřil and Patrice Ossona De Mendez. *Sparsity*. Jan. 2012. DOI: 10.1007/978-3-642-27875-4. URL: <https://doi.org/10.1007/978-3-642-27875-4> (cit. on p. 9).

- [NG04] M. E. J. Newman and M. Girvan. “Finding and evaluating community structure in networks”. In: *Physical Review E* 69.2 (2004), p. 026113. DOI: 10.1103/physreve.69.026113 (cit. on p. 1).
- [RAB09] Martin Rosvall, Daniel Axelsson, and Carl T. Bergstrom. “The map equation”. In: *European Physical Journal-special Topics* 178.1 (Nov. 2009), pp. 13–23. DOI: 10.1140/epjst/e2010-01179-1. URL: <https://doi.org/10.1140/epjst/e2010-01179-1> (cit. on p. 1).
- [TWV19] V. A. Traag, L. Waltman, and N. J. Van Eck. “From Louvain to Leiden: guaranteeing well-connected communities”. In: *Scientific reports* 9.1 (2019), pp. 1–12 (cit. on p. 1).

A Definitions of categorical language